

		40k connections @10k req/s				20k connections @10k req/s					
Erlang version		R15B02				R15B01		R15B02			
CloudI version		1.1.0				1.0.0		1.1.0			
HTTP server		misultin	cowboy	misultin overhead		misultin	change	misultin	cowboy	misultin overhead	
latency mean max (msec) *	erlang	180	19	161	847.37%	210	-2.44%	205	41	164	400.00%
	c	475	395	80	20.25%	38	13.64%	44	34	10	29.41%
	java	1190	1000	190	19.00%	235	32.47%	348	335	13	3.88%
	python_c	1590	1550	40	2.58%	n/a	n/a	620	610	10	1.64%
	python	7950	7500	450	6.00%	3100	16.22%	3700	3400	300	8.82%
	ruby	8950	8000	950	11.88%	3400	16.56%	4075	3600	475	13.19%
latency max sample (msec) *	erlang	650	115	535	465.22%	800	14.89%	940	315	625	198.41%
	c	600	415	185	44.58%	138	13.75%	160	187	-27	-14.44%
	java	1280	1130	150	13.27%	320	27.27%	440	480	-40	-8.33%
	python_c	1725	1715	10	0.58%	n/a	n/a	700	780	-80	-10.26%
	python	7950	7525	425	5.65%	3200	16.88%	3850	3510	340	9.69%
	ruby	8950	8050	900	11.18%	3475	20.11%	4350	3700	650	17.57%
peak throughput sustained (req/s) (red * if not sustained)	erlang	10626.9	10630.2	3.3	0.03%	9633.7	1.07%	9738.2	9863.7	125.5	1.29%
	c	9917.1	10338.4	421.3	4.25%	10172.5	-0.16%	10156.3	10080.8	-75.5	-0.74%
	java	8809.7	9106.6	296.9	3.37%	10060.1	-9.92%	9152.6	9184.5	31.9	0.35%
	python_c	8323.0	8307.3	-15.7	-0.19%	n/a	n/a	8500.8	8450.7	-50.1	-0.59%
	python	3862.3	5400.0	1537.7	39.81%	4423.2	-10.99%	3985.2	4227.5	242.3	6.08%
	ruby	3535.6	3900.0	364.4	10.31%	4098.0	-13.46%	3611.8	4000.0	388.2	10.75%

\* all red data was estimated from the tsung graphs

#### python\_c speedup w/cowboy

	40k	20k
mean	4.8	5.6
max	4.4	4.5